

CSCI 2270

Data Structures and Algorithms

Lecture 3—Memory part 2

Elizabeth White

elizabeth.white@colorado.edu

Office hours: ECCS 112

Wed 9:30am-11:00am

Thurs 10:00am-11:30am

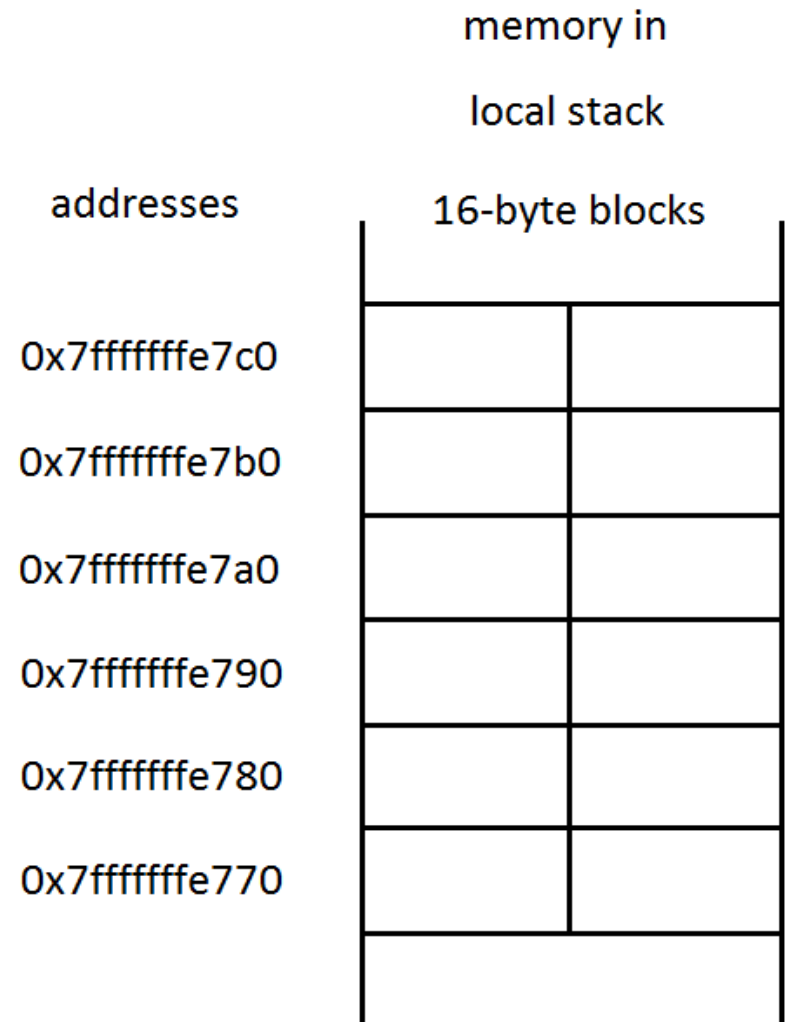
Administrivia

- HW1 will post today
 - It's posting early and will be due Sunday, Feb. 2nd
 - You have not seen everything you'll need for this, so don't jump on it and drive yourself crazy
 - We'll cover function parameters and classes next week
 - LAs will begin to come online today, too
- HW0 is due this week
- No lab next week
- No lecture on Monday
- Lab resumes week after next

Integers at compile time

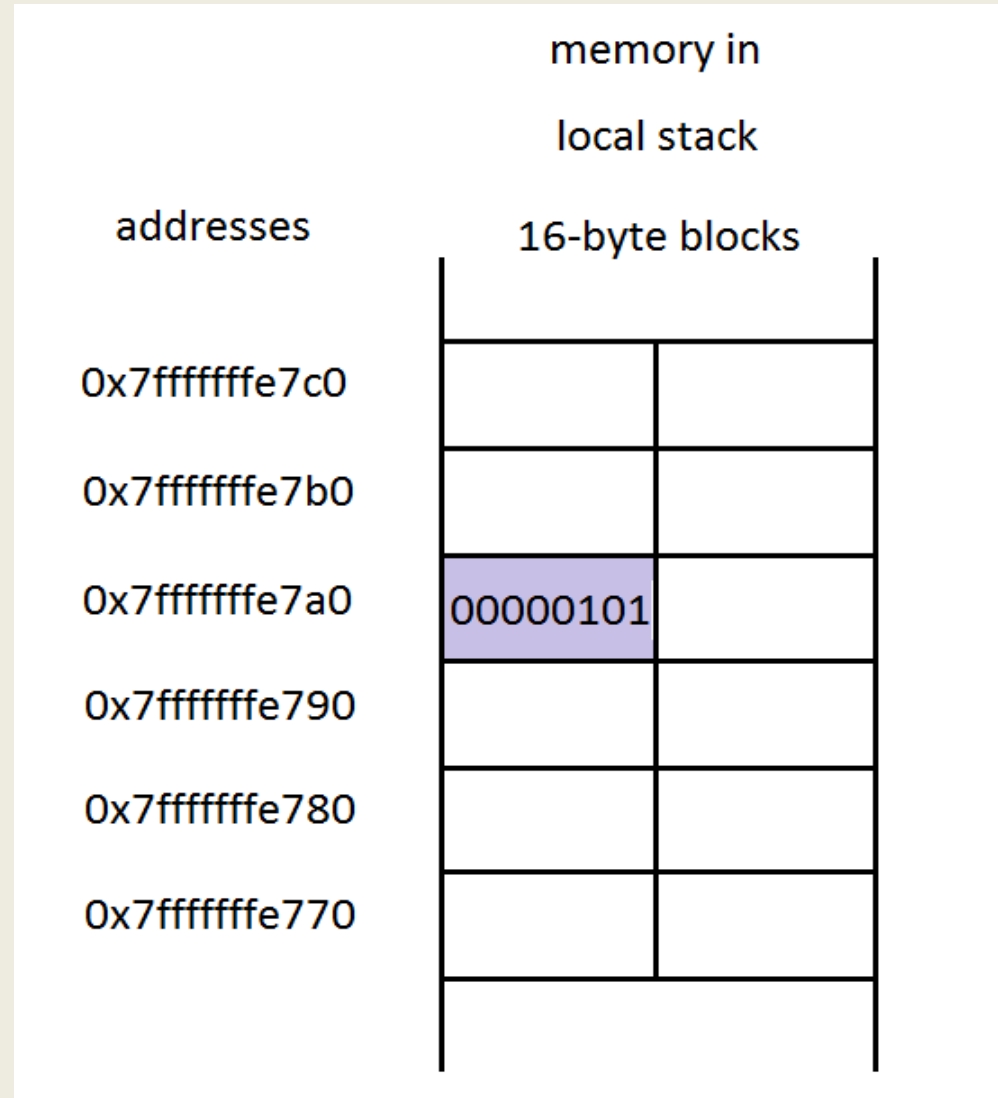
The compiler reserves that memory for a using a local array of memory (called the stack).

I'm drawing it in 8-bit blocks, and I'm showing you the hexadecimal addresses of some of the blocks.



Integers at compile time

After the compiler sees the line `int a = 5;` it reserves a spot for that integer `a` on this stack of memory, and writes the 5 into this spot (in binary). If each slot is a byte, this `int` will take up 4 slots. The operating system makes sure that no other stuff gets written into here, for as long as `a` exists.



Integers at compile time

```
void example1()  
{  
    int a = 5;  
}
```

At the end of our function, just at the closing bracket }, the `int a` gets destroyed. Destroying it means that *we no longer keep that memory reserved* on the stack for `a`.

Now other variables can be written into `a`'s former slot in the memory stack. We can't predict exactly when this will happen, but it's only a matter of time before something else writes data over this slot.

Integers at compile time

```
void example2()  
{  
    int a = 5;  
    cout << a << endl;  
}
```

When it runs, this code prints out the current value of **a** in memory. That's 5.

Integers at compile time

```
void example3()  
{  
    int a = 5;  
    cout << &a << endl;  
}
```

This code prints out the current address where `a` is stored in the stack memory. We get the address of `a`, instead of its value, by putting an ampersand in front (`&a`). Notice that this address is one of those hexadecimal numbers.

From this example, you can deduce that any variable in C++ also knows where it lives in memory.

Integers at compile time

```
void example4()  
{  
    int a = 5;  
    int* a_ptr = &a;  
}
```

This code makes a separate variable called `a_ptr` to store the address of `a`. Again, we get the address by using the `&` sign in front of the variable name.

Integers at compile time

```
void example4()  
{  
    int a = 5;  
    int* a_ptr = &a;  
}
```

This `a_ptr` is a *pointer* to `a`. A pointer stores an address in memory where a variable is living. To make a pointer to an `int`, we need to add the `*` to the `int` type when we declare it, as below.

```
int* a_ptr = &c;
```

It's easy to misplace `&` and `*` at first.

Integers at compile time

```
void example6()  
{  
    int a = 5;  
    int* a_ptr = &a;  
    a = a + 1;  
}
```

We can still change `a` in the normal way, as in the last line above, where it becomes 6.

Integers at compile time

```
void example7()  
{  
    int a = 5;  
    int* a_ptr = &a;  
    *a_ptr = *a_ptr + 1; // a is now 6  
}
```

Here, at the closing bracket, we've used `a_ptr` to change the value of `a` (but not `a_ptr`). Instead of 5, `a` is now 6.

Integers at compile time

```
void example7()  
{  
    int a = 5;  
    int* a_ptr = &a;  
    *a_ptr = *a_ptr + 1; // a is now 6  
}
```

We change `a` by dereferencing `a_ptr`, and we get that by saying `*a_ptr`. Dereferencing means that we find the address that `a_ptr` stores, and then we go to that memory slot and grab `a`, which is stored there. (It's a little extra work.)

Integers at compile time

```
void example8()  
{  
    int a = 5;  
    int* a_ptr = &a;  
    a++;           // a becomes 6  
    (*a_ptr)++;   // a becomes 7  
}
```

This code is similar to examples 6 and 7. First we use the `++` operator to increase `a` by 1, rather than doing a direct addition like `+ 1`. And then we dereference `a_ptr` to get `a` and increase that by 1 again, using the `++` operator. Now you can see how we can change `a`'s value using either its name or its address.

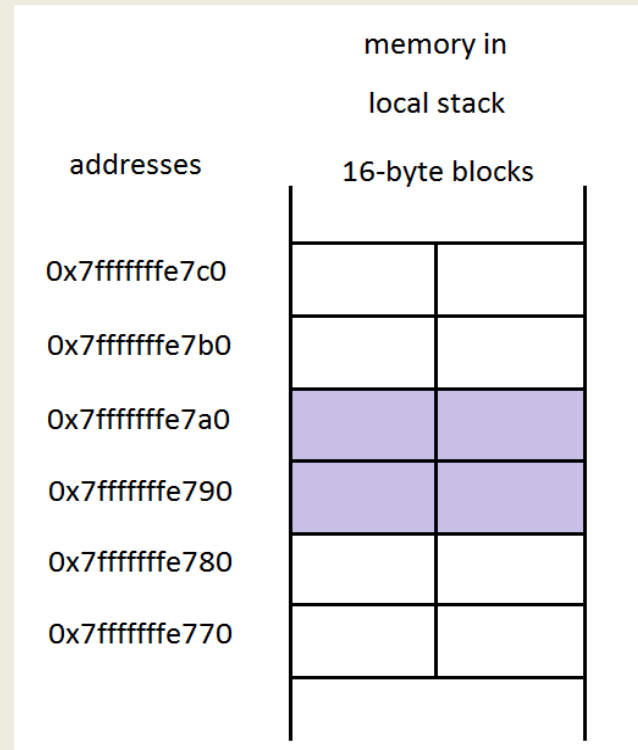
Constant size integer array

```
void example9()  
{  
    int b[4];  
    b[0] = 1;  
    b[1] = 2;  
    b[2] = 4;  
    b[3] = 8;  
}
```

We can make a variable that's an array of integers, instead of one single integer. This code is telling the compiler to find room for 4 integers (that's what the **b[4]** does).

Constant size integer arrays

```
void example9()  
{  
    int b[4];  
    b[0] = 1;  
    b[1] = 2;  
    b[2] = 4;  
    b[3] = 8;  
}
```



For arrays, the compiler stores each integer in adjacent locations in memory. After the line `int b[4];`, it reserves four integer size blocks, one right after the next, for this array.

Constant size integer arrays

```
void example9()  
{  
    int b[4];  
    b[0] = 1;  
    b[1] = 2;  
    b[2] = 4;  
    b[3] = 8;  
}
```

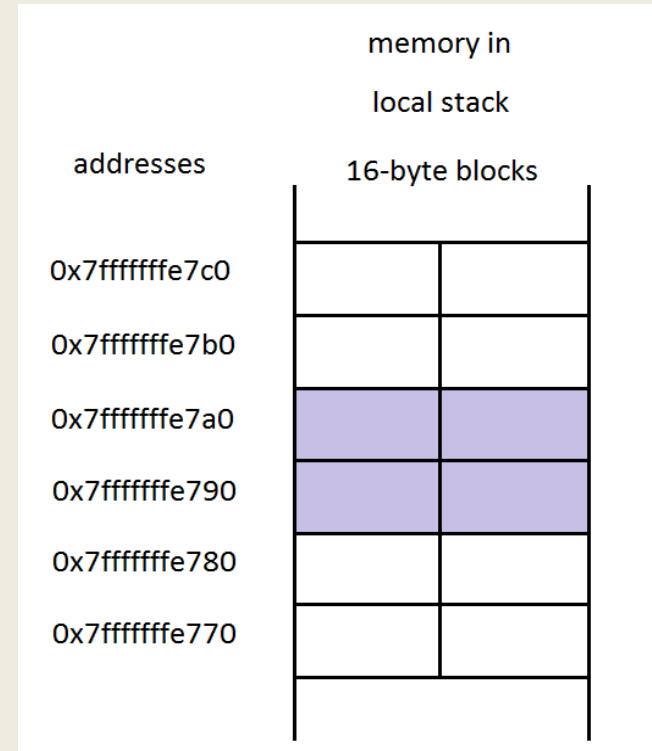
Then we can assign to any of those 4 integers, from `b[0]` to `b[3]`.

Note that arrays in C++ count from 0, not from 1.

Naughty constant size integer array

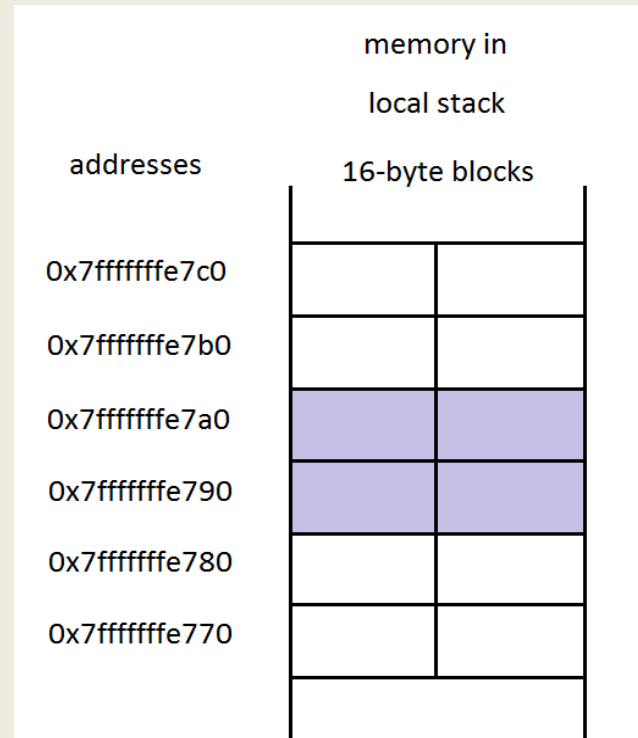
```
void example5()  
{  
    int b[4];  
    b[4] = 5;  
}
```

What happens here?



Why is array size constant?

```
void example9()  
{  
    int b[4];  
    b[0] = 1;  
    b[1] = 2;  
    b[2] = 4;  
    b[3] = 8;  
}
```



An array like `b`, which gets declared with a size in the square brackets (`[]`), is stuck at its starting size forever. We can change the integers in `b[0]` through `b[3]`, but *we can't change `b`'s size* from 4 to something like 10 once we build it with 4 slots.

Floating point numbers

A 4-byte (single precision) floating point number uses:

- 1 bit for the sign,
- 8 bits for the exponent, and
- 23 bits for the binary digits.*

An 8-byte (double precision) floating point number uses:

- 1 bit for the sign,
- 11 bits for the exponent, and
- 52 bits for the binary digits.*

* To a first approximation, anyway. Take 2400 to learn the details. This is all you need for 2270.

Floating point numbers in C++

C++ offers you the choice between floats and doubles.

Floats in the VM

Take up 32 bits of memory, or 4 bytes.

Smallest magnitude: $1.17549\text{e-}38$

Largest magnitude: $3.40282\text{e+}38$

Doubles in the VM

Take up 64 bits of memory, or 8 bytes.

Smallest magnitude: $2.22507\text{e-}308$

Largest magnitude: $1.79769\text{e+}308$

Floating point numbers

Like integers, computers store floating point numbers in binary form. Consider these little numbers:

Decimal	Binary
0	0
0.5	0.1
0.25	0.01
0.125	0.001

For binary numbers,

a 1 in the first place after the decimal (.) is 2^{-1} ,
and a 1 in the next place is 2^{-2} ,
and so forth. All still powers of 2.

Floating point numbers

Some familiar numbers don't work out nicely in binary.

Decimal	Binary
0	0
0.1	0.000110011...
0.333.... (= 1/3)	0.010101...
0.2	0.00110011...

Why should we care?

http://sydney.edu.au/engineering/it/~alum/patriot_bug.html

Time step: 0.1 sec

Accumulated error after 100 hours: 0.34 sec

Error in missile position: 690 m

Boom!